

Evaluating the Effect of Aggregation on Value Function Approximation

COMP 767 – Reinforcement Learning

Mathieu Nassif

McGill University

March 10, 2017

Value Function Approximation

Approximate $v(s)$ using a parameter θ and the function $\hat{v}(s, \theta)$.

Aggregation: We aggregate states into groups and compute an aggregated value for each group.

Effect of Groups

We want to evaluate the effect of the aggregation strategy on

- Error in the limit
- Speed of convergence
- Variance

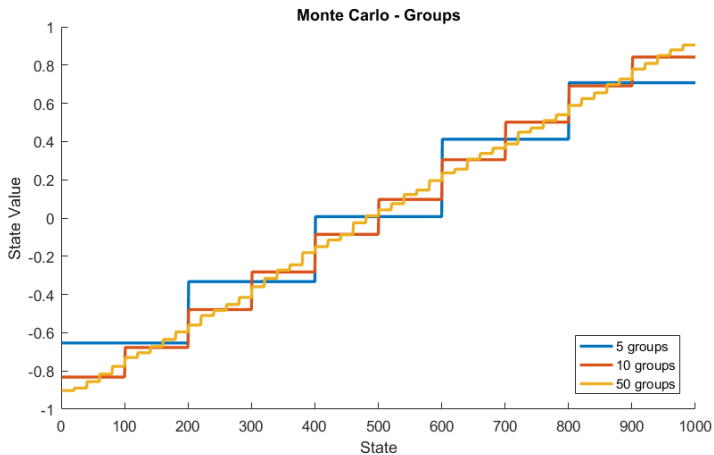
Description of the MDP

- 1000-state random walk
 - 2 actions: left or right
 - Size of jump chosen randomly from 1 to 100 uniformly
 - +1 for reaching right end, -1 for reaching left end, 0 otherwise
 - No discount ($\gamma = 1$)
-
- MC and TD(0) algorithms
 - For all experiments, $\alpha = 0.0004$

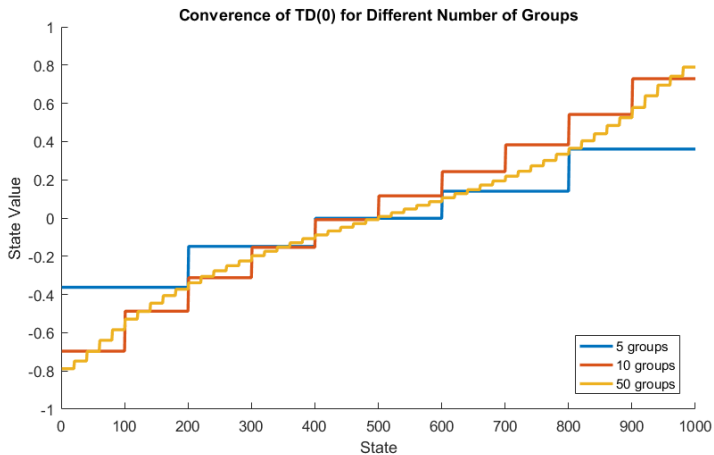
Errors I

- 100 000 episodes
- 5, 10, and 50 groups

Errors II



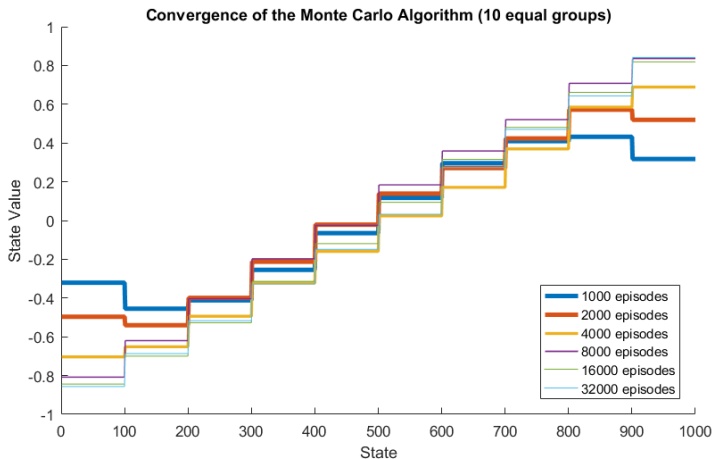
Errors III




Speed of Convergence I

- 32000 episodes (shown after first 1000, 2000, 4000, 8000, 16000, 32000)
- 10 and 50 groups

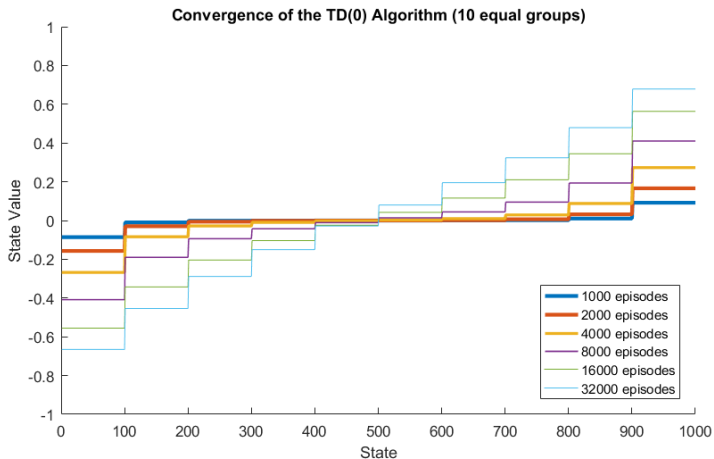
Speed of Convergence II




Speed of Convergence III



Speed of Convergence IV



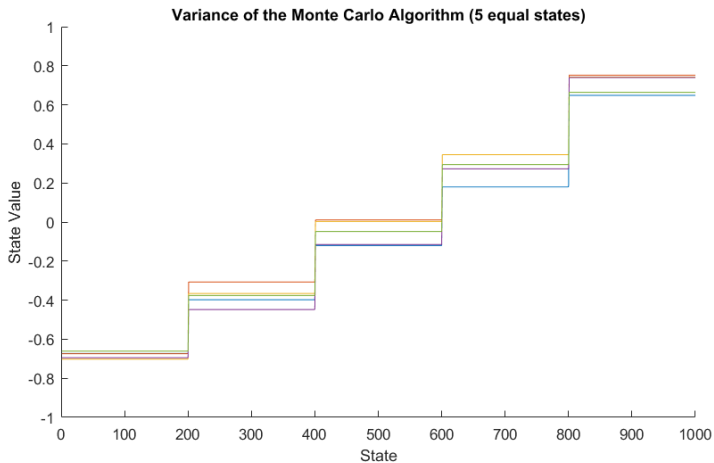
Speed of Convergence V




Variance I

- 5000 episodes
- 5 and 10 groups
- Repeated 5 times each

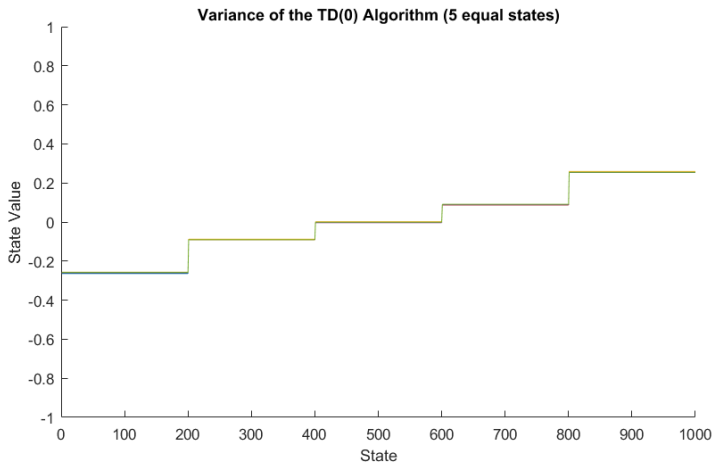
Variance II



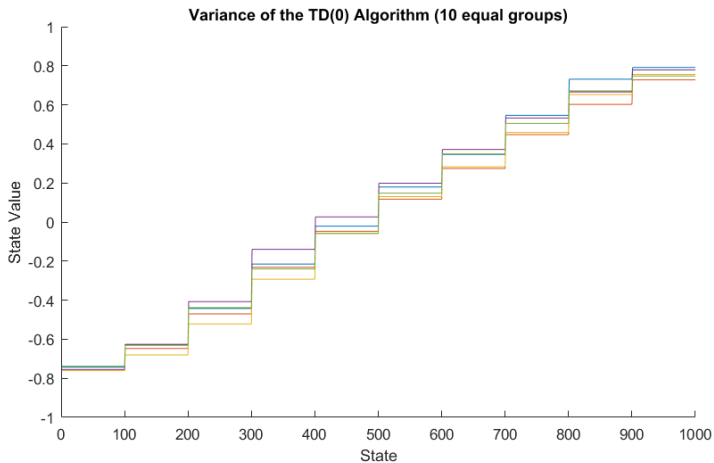
Variance III



Variance IV



Variance V



Difference Grouping Strategies I

Grouping Strategies

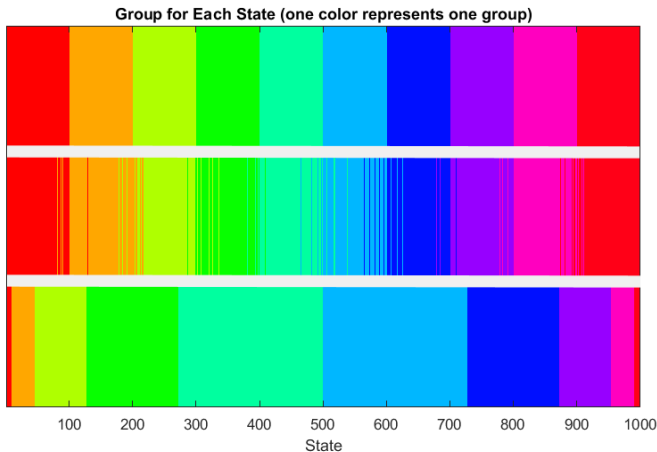
- Constant size
- Constant size with noise

A state near the limit of the group have a non-zero probability of being in the previous/next group.
- Larger groups in the center, thinner near the ends


Difference Grouping Strategies II

Visualization of the groups

(from top to bottom: Constant, Constant with noise, Thinner at the ends)



Difference Grouping Strategies III



Difference Grouping Strategies IV

